

Prinshi Jha

CSE(DS)/23

## Experiment no 4

### Implementation of stochastic gradient descent

```
import numpy as np

def stochastic_gradient_descent(X, y, learning_rate, num_epochs):
    num_samples, num_features = X.shape
    theta = np.zeros(num_features)
    for epoch in range(num_epochs):
        for i in range(num_samples):
            random_index = np.random.randint(num_samples)
            xi = X[random_index:random_index+1]
            yi = y[random_index:random_index+1]
            gradient = np.dot(xi.T, (np.dot(xi, theta) - yi))
            theta -= learning_rate * gradient
    return theta


# Example usage
# Generate some sample data
```

```
np.random.seed(0)  
X = 2 * np.random.rand(100, 3)  
y = 4 + np.dot(X, np.array([3, 2, 1])) + np.random.randn(100)
```

learning\_rate = 0.01

num\_epochs = 100

```
theta = stochastic_gradient_descent(X, y, learning_rate, num_epochs)  
print("Optimal theta:", theta)
```



The screenshot shows a Jupyter Notebook cell with the following content:

```
Untitled11.ipynb  
File Edit View Insert Runtime Tools Help All changes saved  
+ Code + Text  
✓ ⏴ import numpy as np  
{x} def stochastic_gradient_descent(X, y, learning_rate, num_epochs):  
    num_samples, num_features = X.shape  
    theta = np.zeros(num_features)  
  
    for epoch in range(num_epochs):  
        for i in range(num_samples):  
            random_index = np.random.randint(num_samples)  
            xi = X[random_index:random_index+1]  
            yi = y[random_index:random_index+1]  
            gradient = np.dot(xi.T, (np.dot(xi, theta) - yi))  
            theta -= learning_rate * gradient  
  
    return theta  
  
# Example usage  
# Generate some sample data  
np.random.seed(0)  
X = 2 * np.random.rand(100, 3)  
y = 4 + np.dot(X, np.array([3, 2, 1])) + np.random.randn(100)  
  
learning_rate = 0.01  
num_epochs = 100  
  
theta = stochastic_gradient_descent(X, y, learning_rate, num_epochs)  
print("Optimal theta:", theta)
```

The cell has been run, and the output is displayed below the code:

```
Optimal theta: [3.79571426 3.1991076 2.44352241]
```

Implementation of mini batch gradient descent

```
import numpy as np
```

```
# Define the function to be minimized and its gradient
```

```
def f(x):
```

```
    return x**2
```

```
def df(x):
```

```
    return 2*x
```

```
# Define the mini-batch gradient descent function
```

```
def minibatch_gradient_descent(x_init, learning_rate, batch_size, num_iterations):
```

```
    x = x_init
```

```
    for i in range(num_iterations):
```

```
        # Generate random mini-batch samples
```

```
        indices = np.random.choice(len(x), size=batch_size)
```

```
        batch = x[indices]
```

```
        # Compute the gradient of the mini-batch loss
```

```
        gradient = np.mean(df(batch))
```

```
        # Update the parameter using the learning rate and mini-batch gradient
```

```
        x -= learning_rate * gradient
```

```
return x

# Test the mini-batch gradient descent function
x_init = np.array([1.0, 2.0, 3.0, 4.0])
learning_rate = 0.1
batch_size = 2
num_iterations = 100

result = minibatch_gradient_descent(x_init, learning_rate, batch_size,
num_iterations)
print("Result:", result)
```

```
▶ import numpy as np

# Define the function to be minimized and its gradient
def f(x):
    return x**2

def df(x):
    return 2*x

# Define the mini-batch gradient descent function
def minibatch_gradient_descent(x_init, learning_rate, batch_size, num_iterations):
    x = x_init
    for i in range(num_iterations):
        # Generate random mini-batch samples
        indices = np.random.choice(len(x), size=batch_size)
        batch = x[indices]

        # Compute the gradient of the mini-batch loss
        gradient = np.mean(df(batch))

        # Update the parameter using the learning rate and mini-batch gradient
        x -= learning_rate * gradient

    return x

# Test the mini-batch gradient descent function
x_init = np.array([1.0, 2.0, 3.0, 4.0])
learning_rate = 0.1
batch_size = 2
num_iterations = 100

result = minibatch_gradient_descent(x_init, learning_rate, batch_size, num_iterations)
print("Result:", result)
```

▶ Result: [-1.4231927 -0.4231927 0.5768073 1.5768073]